



RTX 5080 · same train inputs and shared heads/loss · 3 timing repetitions

Paired batch processing cost



GPU memory trade-off



32 measured batches after 5 warm-ups; evaluation includes target loss; no accuracy or full-epoch measurement.